

## 8.2 Querying and Merging

**Name:** Calingo, Christian Lei

**Section:** CPE22S3

**Course:** Computational Thinking with Python

**Course Code:** CPE311

```
import pandas as pd
weather = pd.read_csv('data/nyc_weather_2018.csv')
weather.head()
```

	date	datatype	station	attributes	value
0	2018-01-01T00:00:00	PRCP	GHCND:US1CTFR0039	„N,0800	0.0
1	2018-01-01T00:00:00	PRCP	GHCND:US1NJBG0015	„N,1050	0.0
2	2018-01-01T00:00:00	SNOW	GHCND:US1NJBG0015	„N,1050	0.0
3	2018-01-01T00:00:00	PRCP	GHCND:US1NJBG0017	„N,0920	0.0
4	2018-01-01T00:00:00	SNOW	GHCND:US1NJBG0017	„N,0920	0.0

Next steps: [View recommended plots](#)

## Querying DataFrames

The `query()` method is an easier way of filtering based on some criteria. For example, we can use it to find all entries where snow was recorded:

```
snow_data = weather.query('datatype == "SNOW" and value > 0')
snow_data.head()
```

	date	datatype	station	attributes	value
31	2018-01-04T00:00:00	SNOW	GHCND:US1NJBG0015	„N,1600	229.0
34	2018-01-04T00:00:00	SNOW	GHCND:US1NJBG0017	„N,0830	10.0
38	2018-01-04T00:00:00	SNOW	GHCND:US1NJBG0018	„N,0910	46.0
45	2018-01-05T00:00:00	SNOW	GHCND:US1NJBG0017	„N,0720	102.0
49	2018-01-05T00:00:00	SNOW	GHCND:US1NJBG0018	„N,1230	183.0

Using `.loc()` to check if output is the same

```
snow_data2 = weather.loc[(weather["datatype"] == "SNOW") & (weather["value"] > 0)]
snow_data2.head()
```

`#.loc()` can give us the same output however it is easire to use `.query()`

	date	datatype	station	attributes	value
31	2018-01-04T00:00:00	SNOW	GHCND:US1NJBG0015	„N,1600	229.0
34	2018-01-04T00:00:00	SNOW	GHCND:US1NJBG0017	„N,0830	10.0
38	2018-01-04T00:00:00	SNOW	GHCND:US1NJBG0018	„N,0910	46.0
45	2018-01-05T00:00:00	SNOW	GHCND:US1NJBG0017	„N,0720	102.0
49	2018-01-05T00:00:00	SNOW	GHCND:US1NJBG0018	„N,1230	183.0

Equivalent to quering the data/weather.db SQLite database for `SELECT * FROM weather WHERE datatype == "SNOW" AND value > 0`

```
import sqlite3
with sqlite3.connect('data/weather.db') as connection:
    snow_data_from_db = pd.read_sql('SELECT * FROM weather WHERE datatype == "SNOW" AND value > 0', connection)
    snow_data.reset_index().drop(columns='index').equals(snow_data_from_db)
```

True

This is also equivalent to creating Boolean masks:

```
weather[(weather.datatype == 'SNOW') & (weather.value > 0)].equals(snow_data)
```

True

## ✓ Merging DataFrames

We have data for many different stations each day; however, we don't know what the stations are just their IDs. We can join the data in the data/weather\_stations.csv file which contains information from the stations endpoint of the NCEI API. Consult the weather\_data\_collection.ipynb notebook to see how this was collected. It looks like this:

```
station_info = pd.read_csv('data/weather_stations.csv')
station_info.head()
```

	id	name	latitude	longitude	elevation	
0	GHCND:US1CTFR0022	STAMFORD 2.6 SSW, CT US	41.064100	-73.577000	36.6	
1	GHCND:US1CTFR0039	STAMFORD 4.2 S, CT US	41.037788	-73.568176	6.4	
2	GHCND:US1NJBG0001	BERGENFIELD 0.3 SW, NJ US	40.921298	-74.001983	20.1	
3	GHCND:US1NJBG0002	SADDLE BROOK TWP 0.6 E, NJ US	40.902694	-74.083358	16.8	

Next steps: [View recommended plots](#)

```
weather.head()
```

	date	datatype	station	attributes	value	
0	2018-01-01T00:00:00	PRCP	GHCND:US1CTFR0039	„N,0800	0.0	
1	2018-01-01T00:00:00	PRCP	GHCND:US1NJBG0015	„N,1050	0.0	
2	2018-01-01T00:00:00	SNOW	GHCND:US1NJBG0015	„N,1050	0.0	
3	2018-01-01T00:00:00	PRCP	GHCND:US1NJBG0017	„N,0920	0.0	
4	2018-01-01T00:00:00	SNOW	GHCND:US1NJBG0017	„N,0920	0.0	

Next steps: [View recommended plots](#)

We can join our data by matching up the station\_info.id column with the weather.station column. Before doing that though, let's see how many unique values we have:

```
station_info.id.describe()
```

```
count          320
unique          320
top    GHCND:US1CTFR0022
freq           1
Name: id, dtype: object
```

While station\_info has one row per station, the weather dataframe has many entries per station. Notice it also has fewer uniques

```
weather.station.describe()
```

```
#the count is different since I assigned the value of limit as 10 not 1000
```

```

count          3650
unique          14
top    GHCND:US1NJBG0017
freq          576
Name: station, dtype: object

```

**When working with joins, it is important to keep an eye on the row count. Some join types will lead to data loss:**

```
station_info.shape[0], weather.shape[0]
```

```

#once again, the counts are different since I assigned different values on the previous activity
#instead of assigning the limit to 1000, I limited it to 10

```

```
(320, 3650)
```

**We will be doing this often, it makes more sense to write a function:**

```

def get_row_count(*dfs):
    return [df.shape[0] for df in dfs]
get_row_count(station_info, weather)

```

```
[320, 3650]
```

**The map() function is more efficient than list comprehensions. We can couple this with getattr() to grab any attribute for multiple dataframes**

```

def get_info(attr, *dfs):
    return list(map(lambda x: getattr(x, attr), dfs))
get_info('shape', station_info, weather)

```

```
[(320, 5), (3650, 5)]
```

**By default merge() performs an inner join. We simply specify the columns to use for the join. The left dataframe is the one we call merge() on, and the right one is passed in as an argument**

```

inner_join = weather.merge(station_info, left_on='station', right_on='id')
inner_join.sample(5, random_state=0)

# the outputs might be differ in the module
# the reason might be the sudden change in the limit in the previous activity

```

	date	datatype	station	attributes	value	i
3516	2018-09-26T00:00:00	PRCP	GHCND:US1NJBG0010	„N,0800	84.8	GHCND:US1NJBG001
569	2018-05-23T00:00:00	PRCP	GHCND:US1NJBG0015	„N,1120	17.0	GHCND:US1NJBG001

**We can remove the duplication of information in the station and id columns by renaming one of them before the merge and then simply using on**

```

weather.merge(station_info.rename(dict(id='station')), axis=1, on='station').sample(5, random_state=0)

#since id and station have the same values, the id got merged with the station

```

	date	datatype	station	attributes	value	name	latiti
3516	2018-09-26T00:00:00	PRCP	GHCND:US1NJBG0010	„N,0800	84.8	RIVER VALE TWP 1.5 S, NJ US	40.9914
569	2018-05-23T00:00:00	PRCP	GHCND:US1NJBG0015	„N,1120	17.0	NORTH ARLINGTON 0.7 WNW, NJ US	40.7914

We are losing stations that don't have weather observations associated with them, if we don't want to lose these rows, we perform a right or left join instead of the inner join

```
left_join = station_info.merge(weather, left_on='id', right_on='station', how='left')
right_join = weather.merge(station_info, left_on='station', right_on='id', how='right')
right_join.tail()
```

	date	datatype	station	attributes	value	id	name
3951	NaN	NaN	NaN	NaN	NaN	GHCND:USW00054787	FARMINGDALE REPUBLIC AIRPORT, NY US
3952	NaN	NaN	NaN	NaN	NaN	GHCND:USW00094728	NY CITY CENTRAL PARK, NY US

The left and right join as we performed above are equivalent because the side that we kept the rows without matches was the same in both cases

```
left_join.sort_index(axis=1).sort_values(['date', 'station']).reset_index().drop(columns='index').equals(
right_join.sort_index(axis=1).sort_values(['date', 'station']).reset_index().drop(columns='index')
)
```

True

Note we have additional rows in the left and right joins because we kept all the stations that didn't have weather observations

```
get_info('shape', inner_join, left_join, right_join)
```

```
[(3650, 10), (3956, 10), (3956, 10)]
```

If we query the station information for stations that have NY in their name, believing that to be all the stations that record weather data for NYC and perform an outer join, we can see where the mismatches occur:

```
outer_join = weather.merge(
station_info[station_info.name.str.contains('NY')],
left_on='station', right_on='id', how='outer', indicator=True
)
outer_join.sample(4, random_state=0).append(outer_join[outer_join.station.isna()].head(2))
```

<ipython-input-31-81b63e73e04e>:5: FutureWarning: The frame.append method is deprecated

```
outer_join.sample(4, random_state=0).append(outer_join[outer_join.station.isna()].head(2))
```

	date	datatype	station	attributes	value	i
538	2018-05-02T00:00:00	SNOW	GHCND:US1NJBG0015	„N,0815	0.0	Na
526	2018-04-23T00:00:00	SNOW	GHCND:US1NJBG0015	„N,1015	0.0	Na
2215	2018-05-20T00:00:00	PRCP	GHCND:US1NJBG0023	„N,0745	12.4	Na
2872	2018-03-01T00:00:00	SNOW	GHCND:US1NJBG0003	„N,0730	0.0	Na

These joins are equivalent to their SQL counterparts. Below is the inner join. Note that to use equals() you will have to do some manipulation of the dataframes to line them up

```
import sqlite3
with sqlite3.connect('data/weather.db') as connection:
    inner_join_from_db = pd.read_sql('SELECT * FROM weather JOIN stations ON weather.station == stations.id',connection)
inner_join_from_db.shape == inner_join.shape
```

True

## Revisit the dirty data from the previous module

```
dirty_data = pd.read_csv(
    'data/dirty_data.csv', index_col='date'
).drop_duplicates().drop(columns='SNWD')
dirty_data.head()
```

	station	PRCP	SNOW	TMAX	TMIN	TOBS	WESF	inclement_weather
date								
2018-01-01T00:00:00	?	0.0	0.0	5505.0	-40.0	NaN	NaN	Na
2018-01-02T00:00:00	GHCND:USC00280907	0.0	0.0	-8.3	-16.1	-12.2	NaN	Fals
2018-01-03T00:00:00	GHCND:USC00280907	0.0	0.0	-4.4	-13.9	-13.3	NaN	Fals

Next steps: [View recommended plots](#)

## We need to create two dataframes for the join. We will drop some unnecessary columns as well for easier viewing

```
valid_station = dirty_data.query('station != "?"').copy().drop(columns=['WESF', 'station'])
station_with_wesf = dirty_data.query('station == "?"').copy().drop(columns=['station', 'TOBS', 'TMIN', 'TMAX'])
```

*\*Our column for the join is the index in both dataframes, so we must specify left\_index and right\_index \**

```
valid_station.merge(
    station_with_wesf, left_index=True, right_index=True
).query('WESF > 0').head()
```

	PRCP_x	SNOW_x	TMAX	TMIN	TOBS	inclement_weather_x	PRCP_y	SNOW_y	WESF	
date										
2018-01-30T00:00:00	0.0	0.0	6.7	-1.7	-0.6		False	1.5	13.0	1.8
2018-03-08T00:00:00	48.8	NaN	1.1	-0.6	1.1		False	28.4	NaN	28.7
2018-03-13T00:00:00	4.1	51.0	5.6	-3.9	0.0		True	3.0	13.0	3.0

The columns that existed in both dataframes, but didn't form part of the join got suffixes added to their names: `_x` for columns from the left dataframe and `_y` for columns from the right dataframe. We can customize this with the `suffixes` argument:

```
valid_station.merge(
    station_with_wesf, left_index=True, right_index=True, suffixes=('', '_?')
).query('WESF > 0').head()
```

	PRCP	SNOW	TMAX	TMIN	TOBS	inclement_weather	PRCP_?	SNOW_?	WESF	incl
date										
2018-01-30T00:00:00	0.0	0.0	6.7	-1.7	-0.6		False	1.5	13.0	1.8
2018-03-08T00:00:00	48.8	NaN	1.1	-0.6	1.1		False	28.4	NaN	28.7
2018-03-13T00:00:00	4.1	51.0	5.6	-3.9	0.0		True	3.0	13.0	3.0

Since we are joining on the index, an easier way is to use the `join()` method instead of `merge()`. Note that the suffix parameter is now `lsuffix` for the left dataframe's suffix and `rsuffix` for the right one's

```
valid_station.join(station_with_wesf, rsuffix='_?').query('WESF > 0').head()
```

	PRCP	SNOW	TMAX	TMIN	TOBS	inclement_weather	PRCP_?	SNOW_?	WESF	incl
date										
2018-01-30T00:00:00	0.0	0.0	6.7	-1.7	-0.6	False	1.5	13.0	1.8	
2018-03-08T00:00:00	48.8	NaN	1.1	-0.6	1.1	False	28.4	NaN	28.7	
2018-03-13T00:00:00	4.1	51.0	5.6	-3.9	0.0	True	3.0	13.0	3.0	

Joins can be very resource-intensive, so it's a good idea to figure out what type of join you need using set operations before trying the join itself. The pandas set operations are performed on the index, so whichever columns we will be joining on will need to be the index. Let's go back to the weather and station\_info dataframes and set the station ID columns as the index

```
weather.set_index('station', inplace=True)
station_info.set_index('id', inplace=True)
```

The intersection will tell us the stations that are present in both dataframes. The result will be the index when performing an inner join

```
weather.index.intersection(station_info.index)
```

```
Index(['GHCND:US1CTFR0039', 'GHCND:US1NJBG0015', 'GHCND:US1NJBG0017',
      'GHCND:US1NJBG0018', 'GHCND:US1NJBG0023', 'GHCND:US1NJBG0030',
      'GHCND:US1NJBG0039', 'GHCND:US1NJBG0003', 'GHCND:US1NJBG0044',
      'GHCND:US1NJS0018', 'GHCND:US1NJBG0010', 'GHCND:US1NJS0019',
      'GHCND:US1NJS0024', 'GHCND:US1NJBG0037'],
      dtype='object')
```

The set difference will tell us what we lose from each side. When performing an inner join, we lose nothing from the weather dataframe

```
weather.index.difference(station_info.index)
```

```
Index([], dtype='object')
```

We lose 153 stations from the station\_info dataframe, however

```
station_info.index.difference(weather.index)
```

```
Index(['GHCND:US1CTFR0022', 'GHCND:US1NJBG0001', 'GHCND:US1NJBG0002',
      'GHCND:US1NJBG0005', 'GHCND:US1NJBG0006', 'GHCND:US1NJBG0008',
      'GHCND:US1NJBG0011', 'GHCND:US1NJBG0012', 'GHCND:US1NJBG0013',
      'GHCND:US1NJBG0020',
      ...,
      'GHCND:USW00014708', 'GHCND:USW00014732', 'GHCND:USW00014734',
      'GHCND:USW00014786', 'GHCND:USW00054743', 'GHCND:USW00054787',
      'GHCND:USW00094728', 'GHCND:USW00094741', 'GHCND:USW00094745',
      'GHCND:USW00094789'],
      dtype='object', length=306)
```

The symmetric difference will tell us what gets lost from both sides. It is the combination of the set difference in both directions

```
ny_in_name = station_info[station_info.name.str.contains('NY')]
ny_in_name.index.difference(weather.index).shape[0]\
+ weather.index.difference(ny_in_name.index).shape[0]\
== weather.index.symmetric_difference(ny_in_name.index).shape[0]
```

```
True
```

The union will show us everything that will be present after a full outer join. Note that since these are sets (which don't allow duplicates by definition), we must pass unique entries for union:

```
weather.index.unique().union(station_info.index)
```

```
Index(['GHCND:US1CTFR0022', 'GHCND:US1CTFR0039', 'GHCND:US1NJBG0001',
      'GHCND:US1NJBG0002', 'GHCND:US1NJBG0003', 'GHCND:US1NJBG0005',
      'GHCND:US1NJBG0006', 'GHCND:US1NJBG0008', 'GHCND:US1NJBG0010',
      'GHCND:US1NJBG0011',
      ...
      'GHCND:USW00014708', 'GHCND:USW00014732', 'GHCND:USW00014734',
      'GHCND:USW00014786', 'GHCND:USW00054743', 'GHCND:USW00054787',
      'GHCND:USW00094728', 'GHCND:USW00094741', 'GHCND:USW00094745',
      'GHCND:USW00094789'],
      dtype='object', length=320)
```

**Note that the symmetric difference is actually the union of the set differences**

```
ny_in_name = station_info[station_info.name.str.contains('NY')]
ny_in_name.index.difference(weather.index).union(weather.index.difference(ny_in_name.index)).equals(
weather.index.symmetric_difference(ny_in_name.index)
)
```

True

## Conclusion

In this activity, we performed different querying and merging of dataframes. Some of my outputs are not the same in the module because I tried changing the limit:

```
import datetime
from IPython import display # for updating the cell dynamically
current = datetime.date(2018, 1, 1)
end = datetime.date(2019, 1, 1)
results = []
while current < end:
    # update the cell with status information
    display.clear_output(wait=True)
    display.display(f'Gathering data for {str(current)}')
    response = make_request('data',
        {
            'datasetid': 'GHCND', # Global Historical Climatology Network - Daily (GHCND) dataset
            'locationid': 'CITY:US360019', # NYC
            'startdate': current,
            'enddate': current,
            'units': 'metric',
            'limit': 10 # max allowed
        })
    if response.ok:
        # we extend the list instead of appending to avoid getting a nested list
        results.extend(response.json()['results'])
        # update the current date to avoid an infinite loop
        current += datetime.timedelta(days=1)
```

Instead of 1000, I used 10 for faster loading. You can merge different dataframes in lots of ways and also merging columns of a dataframe is also possible. However, some conditions are required in order for it to work.